

CS5050 ADVANCED ALGORITHMS

Fall 2023

Assignment 4: Algorithm Analysis

Due Date: 11:59:59 p.m., *Friday, Nov 10, 2023*

Total Points: 100

1. **(20 points)** Suppose we have a min-heap with n distinct keys that are stored in an array $A[1 \dots n]$ (a min-heap is one that stores the smallest key at its root). Given a value x and an integer k with $1 \leq k \leq n$, design an algorithm to determine whether the k -th smallest key in the heap is smaller than x (so your answer should be “yes” or “no”). The running time of your algorithm should be $O(k)$, independent of the size of the heap.

Remark. If we were to find the k -th smallest key of the heap, denoted by y , then the best way would be to perform k times *deleteMin* operations, which would take $O(k \log n)$ time (or using the selection algorithm, which would take $O(n)$ time). Our above problem, however, is actually a *decision problem*. Namely, you only need to decide whether y is smaller than x , and you do not have to know what the exact value of y is. Hence, the problem is easier and we are able to solve it in a faster way, i.e., $O(k)$ time.

2. **(20 points)** Suppose you are given a balanced binary search tree T of n nodes (as discussed in class, each node v has $v.left$, $v.right$, and $v.key$). We assume that no two nodes of T have the same key. Given a value x , the *successor* of x in T is defined as follows: (1) If x is larger than every key in T , then x does not have a successor; (2) if x is equal to a key in T , then the successor of x is x itself; otherwise, the successor of x is the smallest key of T that is larger than x .

For example, in Figure 1, the successor of 19 is 20, the successor of 48 is 48, and 70 does not have a successor.




Figure 1: A binary search tree.

Design an $O(\log n)$ time algorithm to perform the **successor** operations: Given any value x , your algorithm should return the successor of x in T , and return “NULL” if x does not have a successor in T .

Note: Please give the pseudocode of your algorithm.

3. **(20 points)** Suppose you are given a balanced binary search tree T of n nodes (as discussed in class, each node v has $v.left$, $v.right$, and $v.key$). We assume that no two keys in T are equal. Given a value x , the *rank* operation $rank(x)$ is to return the *rank* of x in T , which is defined to be one plus the number of keys

of T smaller than x . For example, if T has 3 keys smaller than x , then $\text{rank}(x) = 4$. Note that x may or may not be a key in T . For example, in Figure 1, $\text{rank}(16) = 3$, $\text{rank}(21) = 6$, $\text{rank}(25) = 7$, $\text{rank}(26) = 8$. We know that T can support the ordinary *search*, *insert*, and *delete* operations, each in $O(\log n)$ time. You are asked to augment T , such that the *rank* operation, as well as the normal *search*, *insert*, and *delete* operations, all take $O(\log n)$ time each.

You must present: (1) the design of your data structure (i.e., how you augment T); (2) the algorithm for implementing the $\text{rank}(x)$ operation (please give the pseudocode); (3) briefly explain why the normal operations *search*, *insert*, and *delete* can still be performed in $O(\log n)$ time each (you do not need to provide the details of these operations).

4. **(20 points)** This problem is concerned with **range queries** (we have discussed a similar problem in class) on a balanced binary search tree T whose keys are distinct (no two keys in T are equal). The range query is a generalization of the ordinary *search* operation. The **range** of a range query on T is defined by a pair $[x_l, x_r]$, where x_l and x_r are real numbers and $x_l \leq x_r$. Note that x_l and x_r may not be the keys in T .

You already know that T can support the ordinary *search*, *insert*, and *delete* operations, each in $O(\log n)$ time, where n is the number of nodes of T . You are asked to design an algorithm to efficiently perform the *range queries*. That is, in each range query, you are given a range $[x_l, x_r]$, and your algorithm should report all keys x stored in T such that $x_l \leq x \leq x_r$. Your algorithm should run in $O(k + \log n)$ time, where k is the number of keys of T in the range $[x_l, x_r]$. In addition, it is required that all keys in $[x_l, x_r]$ be reported in a *sorted order*. Please give the pseudocode for your algorithm.

Remark. Such an algorithm of $O(k + \log n)$ time is an *output-sensitive* algorithm because the running time (i.e., $O(k + \log n)$) is a function of the output size k . As an application of the range queries, suppose the keys of T are student scores in an exam. A range query like $[70, 80]$ would report all scores in the range in sorted order.

5. **(20 points)** Consider one more operation on the above balanced binary search tree T in Problem 4: *range-sum* (x_l, x_r) . Given any range $[x_l, x_r]$ with $x_l \leq x_r$, the operation *range-sum* (x_l, x_r) computes the *sum* of the keys in T that are in the range $[x_l, x_r]$.

You are asked to augment the binary search tree T , such that the *range-sum* (x_l, x_r) operations, as well as the ordinary *search*, *insert*, and *delete* operations, all take $O(\log n)$ time each.

You must present: (1) the design of your data structure (i.e., how you augment T); (2) the algorithm for implementing the *range-sum* (x_l, x_r) operation (please give the pseudocode); (3) briefly explain why the ordinary operations *search*, *insert*, and *delete* can still be performed in $O(\log n)$ time each (you do not need to provide the details of these operations).